

IMPACT pSIDDHI 3.0

Final-Term Submission Document

Common template for all tracks (Custom · Data · Platform) and all semesters, including pSiddhi Lite | Covers the ENTIRE programme, Week 4 through Week 17

READ BEFORE FILLING

- This document records what you have ACTUALLY built and verified across the FULL programme (Week 4–17), measured against your L&D-APPROVED proposal, including any changes formally agreed with L&D during Phase 2. It is an end-to-end record, not a delta on your Mid-Term submission.
- This document is evaluated TOGETHER with your Mid-Term submission document — quote your Mid-Term filename in Section 1 so the two are paired automatically.
- Every deliverable your approved proposal committed to across the full programme appears in Section 3 (D-01 to D-12). Every deliverable marked 'Done' points to verified Evidence IDs in Section 4.
- All 12 sections are strictly preserved for automated AI scoring engine evaluation.

1. Participant & Project Identification

Topic ID (as finalised by L&D)	S4-I-21
Topic Title	Unified Sales Operations System
Participant Name	Hindujashiri Gopu
Employee ID	P403
Track	☐ Custom ☒ Data ☐ Platform
Semester & Category	Semester 4 — Integration Mastery (Capstone)
Participation Type	☒ Regular ☐ pSiddhi Lite
Approved Budget Ceiling	₹2,500 (fixed)
Mid-Term document filename (as uploaded to Moodle)	S4-I-21_Hindujashiri_MidTermDoc.docx
Mid-Term result / feedback received	☒ On track ☐ At risk — actions were assigned ☐ Other
Final Review Window	Week 17

2. Approved Proposal Recap

2.1 Problem Statement (as approved)

The sales operations function across modern B2B organizations suffers from four interlocking structural failures: (1) CRM data integrity fails because reps update deal stages, values, and close dates sporadically, forcing sales managers to make resourcing decisions against stale pipelines; (2) Traditional stage-weighted forecasting applies static percentage multipliers to pipeline stages regardless of rep win rate or deal velocity, producing inaccurate projections; (3) Leading account health signals (support ticket spikes, dropping engagement, stalling expansion) remain invisible until accounts are already churning; and (4) The complete absence of unified AI/ML intelligence leaves deal prioritization and rep coaching to subjective intuition. SalesIQ unifies these disconnected domains into an integrated, ML-driven sales operations intelligence platform.

2.2 Proposed Solution Summary (as approved)

SalesIQ is architected as an enterprise-grade, four-tier unified sales operations system: Layer 1 is a high-throughput Python FastAPI integration engine ingesting three core CRM domains (pipeline opportunities, account health, and sales reps) with Pydantic schema validation into a normalized SQLite persistence layer. Layer 2 is a modern, 4-screen React portal featuring Pipeline Overview, Account Detail, Rep Performance, and AI Intelligence. Layer 3 is an embedded ML analytics engine running three scikit-learn models: a Deal Win Probability Scorer (classification), a Revenue Forecaster (regression), and an Account Health Classifier (classification), all tracked in MLflow. Layer 4 is an AI Sales Intelligence engine synthesizing cross-domain features into grounded executive narratives across three operational scenarios (Strong Quarter, At-Risk Quarter, and Recovery Scenario) with dual Gemini 2.5 Flash and deterministic fallback support.

2.3 Core Tools & AI Components (as approved)

- Data Integration & Persistence: Python 3.12/3.13, FastAPI, SQLite with WAL mode, JSON Server (mock CRM).
- Machine Learning Engine: scikit-learn (auto-selection pipelines), MLflow (experiment & parameter tracking).
- AI Intelligence Engine: Google Gemini 2.5 Flash API + Grounded Deterministic Offline Narrative Engine.
- Frontend Portal: React 18, Vite 5, Recharts (responsive analytics visualizations), Vanilla CSS.
- Quality Assurance & CI: pytest, pytest-cov (≥80% mandatory coverage target), GitHub Actions CI.
- Cloud Infrastructure: Azure Static Web Apps (Free tier) + Azure App Service Linux F1 within the ₹2,500 budget ceiling.

3. Progress Against Approved Plan (Full Programme: Week 4–17)

ID	Planned Deliverable (per approved proposal)	Planned Window	Carried from Mid-Term? (Y/N)	Status	Evidence ID(s)
D-01	Mock CRM (JSON Server) + unified schema + FastAPI scaffold	Week 4	Y	Done	EV-01, EV-02
D-02	FastAPI integration — pipeline + rep domains ingested to SQLite	Week 5	Y	Done	EV-02
D-03	Account-health domain; all 3 domains normalised & queryable	Week 6	Y	Done	EV-02
D-04	React portal — Pipeline Overview + Account Detail screens	Week 7	Y	Done	EV-04, EV-05, EV-06
D-05	ML models — Revenue Forecaster + Win	Week 8	Y	Done	EV-02, EV-08

	Probability Scorer + MLflow				
D-06	Gemini connected; initial AI narrative (Strong Quarter scenario)	Week 9	Y	Done	EV-07
D-07	QA embedded — unit/integration/API/ML/AI tests, ≥80% coverage	Weeks 4–9	Y	Done	EV-03
D-08	Account Health Classifier (3rd ML model, F1 > 0.75) + MLflow tracking	Week 11	N	Done	EV-02, EV-09
D-09	Rep Performance portal screen with quota attainment & leaderboards	Week 11	N	Done	EV-10, EV-11
D-10	AI Intelligence portal screen with interactive 3-scenario selector	Weeks 11–12	N	Done	EV-12, EV-13
D-11	All 3 AI Scenarios live (Strong Quarter, At-Risk Quarter, Recovery)	Week 12	N	Done	EV-12, EV-13
D-12	Azure deployment manifests & GitHub Actions CI green (94% coverage)	Weeks 13–16	N	Done	EV-14, EV-15

3.1 Overall Final Self-Assessment

RFP-defined final checkpoint / definition of done (summarise in 2–3 lines)	Full delivery of a unified sales operations system integrating 3 CRM domains, 4 functional portal screens (< 3s load), 3 trained ML models meeting accuracy targets (MAPE < 15%, Win AUC > 0.75, Health F1 > 0.75), 3 AI narrative scenarios, and ≥80% automated test coverage within the ₹2,500 budget ceiling.
% of overall project completed (honest estimate)	100%
% reported at Mid-Term (from your Mid-Term Section 3.1)	~100% of Phase 1 (Week 4–9 scope)
Is the final solution demonstrable live, end-to-end, at the review?	☒ Yes, end-to-end ☐ Yes, partially ☐ No, recording/screenshots only

4. Evidence Pack (Full Programme: Week 4–17)

4.1 Evidence Index

Evidence ID	Caption — what does this prove?	Deliverable ID(s)	Verifiable link (if any)	Carried from Mid-Term?
-------------	---------------------------------	-------------------	--------------------------	------------------------

EV-01	Synthetic CRM data generation: 660 records generated, 15 invalid rejected	D-01	data/generate_data.py	Yes — same as Mid-Term EV-01
EV-02	End-to-end pipeline CLI: Ingest 3 domains, train 3 models, score deals & accounts	D-02, D-03, D-05, D-08	backend/pipeline_cli.py	No (progressed with 3rd ML model)
EV-03	Test suite execution: 34 tests passing with 94% measured coverage	D-07, D-12	backend/tests/	No (progressed from 28 to 34 tests)
EV-04	FastAPI live Swagger documentation & pipeline overview endpoint	D-04	http://localhost:8000/docs	Yes — same as Mid-Term EV-04
EV-05	Portal Screen 1 — Pipeline Overview (forecast bar chart & top deals)	D-04	http://localhost:5173/	Yes — same as Mid-Term EV-05
EV-06	Portal Screen 2 — Account Detail (health mix & renewal tracker)	D-03, D-04	http://localhost:5173/accounts	Yes — same as Mid-Term EV-06
EV-07	AI Intelligence narrative generation (Strong Quarter scenario)	D-06	http://localhost:8000/api/intelligence/narrative	Yes — same as Mid-Term EV-07
EV-08	MLflow experiment tracking: parameters, metrics & model registry	D-05, D-08	backend/mlruns/	Yes — same as Mid-Term EV-08
EV-09	Account Health Classifier (3rd model) cross-validation & F1 = 0.769 > 0.75	D-08	backend/app/ml/train.py	No (new Phase 2)
EV-10	Portal Screen 3 — Rep Performance: Team KPIs, Quota Attainment & Velocity chart	D-09	http://localhost:5173/rep	No (new Phase 2)
EV-11	Portal Screen 3 — Rep Leaderboard: 40 reps, quotas, attainment & coaching badges	D-09	http://localhost:5173/rep	No (new Phase 2)
EV-12	Portal Screen 4 — AI Sales Intelligence: 3-scenario cockpit & grounded narrative	D-10, D-11	http://localhost:5173/intelligence	No (new Phase 2)
EV-13	Portal Screen 4 — AI Intelligence:	D-10, D-11	http://localhost:5173/intelligence	No (new Phase 2)

	Priority Focus Deals & Account Churn Watchlist			2)
EV-14	Azure Deployment manifests (staticwebapp.config.json, Dockerfile, startup.sh)	D-12	docs/AZURE_DEPLOYMENT.md	No (new Phase 2)
EV-15	GitHub Actions CI green workflow enforcing pytest coverage on main branch	D-12	.github/workflows/ci.yml	No (new Phase 2)

4.2 Evidence Blocks (paste screenshots here)

EV-01 — Synthetic CRM Data Generation & Schema Validation

What this proves	Synthetic generation of 660 CRM records across 3 domains (500 deals, 120 accounts, 40 reps) with 15 injected invalid records correctly rejected at ingestion.
Deliverable ID (Sec 3)	D-01
Date of Development / Testing	10-Jun-2026 (Week 4)
Verifiable link	data/generate_data.py
Carried from Mid-Term?	☒ Yes — same as Mid-Term EV-01 ☐ No (new / progressed)

[Screenshot referenced from Mid-Term document or repository verification artifacts]

EV-02 — End-to-End Pipeline CLI (Ingestion, 3 ML Models Training & Scoring)

What this proves	Unified execution of 3-domain ingestion, cross-validation model training for all 3 models (Win Scorer AUC 0.9643, Rev Forecaster MAPE 0.0537, Health Classifier F1 0.7690), deal scoring, and account health classification.
Deliverable ID (Sec 3)	D-02, D-03, D-05, D-08
Date of Development / Testing	03-Sep-2026 (Week 16)
Verifiable link	backend/pipeline_cli.py
Carried from Mid-Term?	☐ Yes ☒ No (new / progressed — upgraded with 3rd ML model)

[Screenshot referenced from Mid-Term document or repository verification artifacts]

EV-03 — Automated QA Test Suite (34 Tests Passing, 94% Measured Coverage)

What this proves	Complete automated test coverage across schema, integration, ML, API, and AI layers. 34 tests passing in 52s with 94% coverage, exceeding the ≥80% mandate.
-------------------------	---

Deliverable ID (Sec 3)	D-07, D-12
Date of Development / Testing	03-Sep-2026 (Week 16)
Verifiable link	backend/tests/
Carried from Mid-Term?	☐ Yes ☒ No (new / progressed — expanded from 28 to 34 tests)

The screenshot shows an IDE window titled "SalesIQ - Copy - AntigraVity IDE - ci.yml". The editor displays a GitHub workflow file with the following content:

```

9 jobs:
10   backend:
11
21     python-version: "3.12"
22     cache: pip
23
24     - name: Install dependencies
25       run: pip install -r requirements.txt
26

```

The terminal output shows the execution of a pytest command in a PowerShell environment. The output includes a deprecation warning about the "asyncio_default_fixture_loop_scope" option, followed by a successful test session summary:

```

PS C:\Users\hindujashiri.gopu\Documents\SalesIQ - Copy> .\backend\.venv\Scripts\pytest.exe backend\tests -v
C:\Users\hindujashiri.gopu\Documents\SalesIQ\backend\.venv\Lib\site-packages\pytest_asyncio\plugin.py:207: PytestDeprecationWarning: The configuration option "asyncio_default_fixture_loop_scope" is unset.
The event loop scope for asynchronous fixtures will default to the fixture caching scope. Future versions of pytest-asyncio will default the loop scope for asynchronous fixtures to function scope. Set the default fixture loop scope explicitly in order to avoid unexpected behavior in the future. Valid fixture loop scopes are: "function", "class", "module", "package", "session"


  warnings.warn(PytestDeprecationWarning(_DEFAULT_FIXTURE_LOOP_SCOPE_UNSET))
===== test session starts =====
platform win32 -- Python 3.13.14, pytest-8.3.4, pluggy-1.6.0
rootdir: C:\Users\hindujashiri.gopu\Documents\SalesIQ - Copy\backend
configfile: pytest.ini
plugins: anyio-4.14.2, Faker-33.1.0, asyncio-0.25.0, cov-6.0.0
asyncio: mode=Mode.STRICT, asyncio_default_fixture_loop_scope=None
collected 34 items

backend\tests\test_ai_narrative.py ..... [ 17%]
. [ 20%]
backend\tests\test_api_endpoints.py ..... [ 35%]
... [ 44%]

```

EV-04 — FastAPI OpenAPI / Swagger Documentation & Pipeline Endpoints

What this proves	FastAPI REST API serving live endpoints for pipeline, accounts, reps, and AI intelligence with typed Pydantic response schemas and < 100ms response time.
Deliverable ID (Sec 3)	D-04
Date of Development / Testing	10-Jul-2026 (Week 8)
Verifiable link	http://localhost:8000/docs
Carried from Mid-Term?	☒ Yes — same as Mid-Term EV-04 ☐ No (new / progressed)



EV-05 — Portal Screen 1: Pipeline Overview Dashboard

What this proves	Working Pipeline Overview screen showing KPI summary cards (272 open deals, ₹13.13 Cr pipeline), forecast bar chart comparing stage-weighted vs ML forecast, and top deals table.
Deliverable ID (Sec 3)	D-04
Date of Development / Testing	12-Jul-2026 (Week 9)
Verifiable link	http://localhost:5173/
Carried from Mid-Term?	☒ Yes — same as Mid-Term EV-05 ☐ No (new / progressed)

[Screenshot referenced from Mid-Term document or repository verification artifacts]

EV-06 — Portal Screen 2: Account Detail & Health Distribution

What this proves	Working Account Detail screen rendering portfolio health distribution (Healthy, At-Risk, Critical), ARR, support tickets, and linked deals with renewal risk indicators.
Deliverable ID (Sec 3)	D-03, D-04
Date of Development / Testing	12-Jul-2026 (Week 9)
Verifiable link	http://localhost:5173/accounts
Carried from Mid-Term?	☒ Yes — same as Mid-Term EV-06 ☐ No (new / progressed)

[Screenshot referenced from Mid-Term document or repository verification artifacts]

EV-07 — AI Intelligence Narrative Endpoint (Strong Quarter Scenario)

What this proves	AI executive briefing generated for the Strong Quarter scenario citing factual numbers from the ML pipeline (₹7.31 Cr ML forecast vs ₹5.59 Cr stage baseline).
Deliverable ID (Sec 3)	D-06
Date of Development / Testing	14-Jul-2026 (Week 9)
Verifiable link	http://localhost:8000/api/intelligence/narrative
Carried from Mid-Term?	☒ Yes — same as Mid-Term EV-07 ☐ No (new / progressed)

[Screenshot referenced from Mid-Term document or repository verification artifacts]

EV-08 — MLflow Experiment Tracking & Metric Artifacts

What this proves	MLflow tracking server logging hyperparameters, run metadata, cross-validation scores, and model binaries in backend/mlruns/.
Deliverable ID (Sec 3)	D-05, D-08
Date of Development / Testing	13-Jul-2026 (Week 9)
Verifiable link	backend/mlruns/
Carried from Mid-Term?	☒ Yes — same as Mid-Term EV-08 ☐ No (new / progressed)

[Screenshot referenced from Mid-Term document or repository verification artifacts]

EV-09 — 3rd ML Model: Account Health Classifier Training & F1 > 0.75

What this proves	Auto-selection training comparing DecisionTree, RandomForest, and GradientBoosting on engineered signals (ticket intensity, disengagement, ARR, days to renewal). Achieved F1 = 0.769 – 0.991, beating > 0.75 target.
Deliverable ID (Sec 3)	D-08
Date of Development / Testing	03-Sep-2026 (Week 16)
Verifiable link	backend/app/ml/train.py
Carried from Mid-Term?	☐ Yes ☒ No (new / progressed — delivered Phase 2)

[Screenshot referenced from Mid-Term document or repository verification artifacts]

EV-10 — Portal Screen 3: Rep Performance Analytics & Visualizations

What this proves	Functional 3rd portal screen with 4 team KPIs (Active Reps, 70% Attainment, 59% Win
-------------------------	---

	Rate, 96d Cycle), Quota Attainment bar chart with 100% baseline, and Win Rate vs Cycle velocity chart.
Deliverable ID (Sec 3)	D-09
Date of Development / Testing	03-Sep-2026 (Week 16)
Verifiable link	http://localhost:5173/rep
Carried from Mid-Term?	☐ Yes ☒ No (new / progressed — delivered Phase 2)



EV-11 — Portal Screen 3: Rep Leaderboard with Coaching Badges


What this proves	Full rep roster (40 reps) with interactive region/segment filters, attainment progress bars, and automated coaching badges (President's Club, On Track, Needs Coaching).
Deliverable ID (Sec 3)	D-09
Date of Development / Testing	03-Sep-2026 (Week 16)
Verifiable link	http://localhost:5173/rep
Carried from Mid-Term?	☐ Yes ☒ No (new / progressed — delivered Phase 2)

The screenshot displays a web application interface for sales intelligence. It features a list of sales representatives, each with a profile picture, name, ID, region, company type, revenue, and various performance metrics including touchpoints, days, and coaching needs. The interface is dark-themed and includes a sidebar with navigation options.

ID	Name	Region	Company	Revenue	Touchpoints	Days	Coaching Needs
R008	Crystal Benitez	North	SMB	\$9,50,000	58%	43%	69 touches, 149 days, Needs Coaching (54%)
R007	Alex Chapman	South	SMB	\$14,50,000	62%	44%	17 touches, 95 days, Needs Coaching (82%)
R029	John Hunter MD	South	Mid-Market	\$11,00,000	76%	44%	23 touches, 77 days, On Track (76%)
R040	Teresa Rice	North	Mid-Market	\$13,50,000	70%	40%	66 touches, 91 days, Needs Coaching (79%)
R011	Raymond Valenzuela	South	SMB	\$11,50,000	57%	37%	32 touches, 107 days, Needs Coaching (57%)
R031	Michael Jones	North	Mid-Market	\$12,50,000	65%	36%	73 touches, 156 days, Needs Coaching (83%)
R003	Robert Solis	East	Enterprise	\$23,50,000	11%	33%	120 touches, 97 days, Needs Coaching (31%)
R009	Charles Harris	South	SMB	\$12,00,000	43%	30%	26 touches, 72 days, Needs Coaching (43%)
R010	John Collins	South	Enterprise	\$29,50,000	18%	27%	105 touches, 90 days, Needs Coaching (38%)
R002	Sarah Payne	North	Enterprise	\$17,00,000	51%	25%	34 touches, 70 days, Needs Coaching (51%)
R015	Noah Wagner	West	Mid-Market	\$25,50,000	32%	24%	67 touches, 114 days, Needs Coaching (32%)
R039	Sandy Walker	East	Mid-Market	\$18,50,000	16%	23%	102 touches, 140 days, Needs Coaching (36%)
R004	Dalton Oliver	South	Mid-Market	\$11,50,000	50%	22%	27 touches, 25 days, Needs Coaching (30%)
R017	Jeffrey Barnes	South	Enterprise	\$20,00,000	51%	21%	136 touches, 125 days, Needs Coaching (31%)


EV-12 — Portal Screen 4: AI Sales Intelligence 3-Scenario Executive Cockpit

What this proves	Interactive 4th portal screen featuring tab switcher between Strong Quarter, At-Risk Quarter, and Recovery Scenarios with live pipeline context cards and grounded executive brief.
Deliverable ID (Sec 3)	D-10, D-11
Date of Development / Testing	03-Sep-2026 (Week 16)
Verifiable link	http://localhost:5173/intelligence
Carried from Mid-Term?	☐ Yes ☒ No (new / progressed — delivered Phase 2)



EV-13 — Portal Screen 4: Priority Deals & Account Churn Watchlist

What this proves	Actionable decision support tables ranking top high-expected-value opportunities and flagging high-risk accounts approaching renewal for immediate intervention.
Deliverable ID (Sec 3)	D-10, D-11
Date of Development / Testing	03-Sep-2026 (Week 16)
Verifiable link	http://localhost:5173/intelligence
Carried from Mid-Term?	☐ Yes ☒ No (new / progressed — delivered Phase 2)



EV-14 — Azure Cloud Deployment Configurations & Cost Control

What this proves	Production deployment architecture with staticwebapp.config.json (Azure Static Web Apps) and Dockerfile + startup.sh (Azure App Service Linux) verifying spend fits within ₹2,500 budget (₹1,300 estimated).
Deliverable ID (Sec 3)	D-12
Date of Development / Testing	03-Sep-2026 (Week 16)
Verifiable link	docs/AZURE_DEPLOYMENT.md
Carried from Mid-Term?	☐ Yes ☒ No (new / progressed — delivered Phase 2)

[Screenshot referenced from Mid-Term document or repository verification artifacts]

EV-15 — Automated CI/CD Pipeline on GitHub Actions

What this proves	Automated workflow in .github/workflows/ci.yml executing synthetic CRM data generation, pytest test suite with --cov-fail-under=80 enforcement, and Vite frontend production build.
Deliverable ID (Sec 3)	D-12
Date of Development / Testing	03-Sep-2026 (Week 16)
Verifiable link	https://github.com/HindujaShiriG/pSiddhiSalesQ/actions
Carried from Mid-Term?	☐ Yes ☒ No (new / progressed — delivered Phase 2)

[Screenshot referenced from Mid-Term document or repository verification artifacts]

5. Working Demo, Repository & Live Walkthrough Plan

Code repository URL (GitHub/GitLab)	https://github.com/HindujaShiriG/pSiddhiSalesQ
Final commit ID + date (as of final submission)	b219ea6 (03-Sep-2026)
Deployed / hosted URL (if any)	http://localhost:5173 (Live interactive demo) / Azure Static Web Apps
Notebook / dashboard / other artefact links	FastAPI Swagger Docs: http://localhost:8000/docs MLflow tracking: backend/mlruns/ Coverage Report: backend/htmlcov/index.html

5.1 Live Code Walkthrough Plan

Module / flow you will walk through (preferably AI component)	AI Sales Intelligence Engine & ML Analytics Pipeline: Cross-domain signal extraction from SQLite → 3 scikit-learn models inference (Win Probability, Revenue Forecast, Account Health) → Grounded executive narrative generation across 3 scenarios (Strong Quarter, At-Risk Quarter, Recovery Scenario) → Live React portal rendering with interactive tabs.
Repo path(s) / file(s) for that module	backend/app/ai/narrative.py backend/app/ml/train.py backend/app/ml/features.py frontend/src/screens/AllIntelligence.jsx
Branch to use during the walkthrough	main
Anything the panel should open in advance	None. Complete offline deterministic engine ensures 100% demo reliability without external API keys or cloud dependencies.

6. QA Progress (Full Programme)

Test Type (per approved QA strategy)	Tests written / run (total)	Coverage achieved (measured)	Target (per proposal)	Evidence ID(s)
Unit Tests (Pydantic schema validation & data generation)	6	100% on schemas	≥80%	EV-03
Integration Tests (3-domain CRM ingestion & referential integrity)	5	83% on crm_client	All 3 domains	EV-02, EV-03
API Endpoint Tests (FastAPI TestClient on all 4 portal routes)	8	96–100% on routers	All 4 screens	EV-03, EV-04
ML Accuracy Tests (Win Scorer, Rev Forecaster, Health Classifier)	8	AUC=0.964, MAPE=5.37%, F1=0.769	AUC>0.75, MAPE<15%, F1>0.75	EV-02, EV-09

AI Grounding & Scenario Quality Tests (All 3 scenarios)	7	84% on narrative.py	All 3 scenarios	EV-03, EV-12, EV-13
TOTAL AUTOMATED TEST SUITE	34 passed in 52s	94% Total Coverage	≥80% Mandate	EV-03

7. Tool & Budget Reconciliation (Full Programme)

Tool / Service (approved)	Approved tier & cost	Used in final solution?	Actual cost (₹)	Reason if changed / not used
Azure Backend Hosting (App Service)	Paid · ₹800	☐ Yes ☒ No ☐ Partial	₹0	Deployment container manifests configured; ₹0 incurred during local development & evaluation.
Azure App Service F1 (Linux)	Free · ₹0	☐ Yes ☒ No ☒ Partial	₹0	Production Dockerfile and startup.sh authored and tested locally.
Azure Static Web Apps	Free · ₹0	☐ Yes ☒ No ☒ Partial	₹0	staticwebapp.config.json authored for SPA routing and API proxying.
Python FastAPI + SQLite	Open source · ₹0	☒ Yes ☐ No ☐ Partial	₹0	Core integration backend, REST endpoints, and relational persistence.
JSON Server	Open source · ₹0	☒ Yes ☐ No ☐ Partial	₹0	Mock CRM providing 3 REST domains with file fallback.
scikit-learn + MLflow	Free · ₹0	☒ Yes ☐ No ☐ Partial	₹0	Trained 3 ML models using cross-validation; tracked metrics in MLflow.
Google Gemini 2.5 Flash	Free tier · ₹0	☒ Yes ☐ No ☒ Partial	₹0	Configured with deterministic offline fallback for 100% demo uptime.
React 18 + Recharts	Free · ₹0	☒ Yes ☐ No ☐ Partial	₹0	All 4 portal screens built with responsive analytics and dark theme.
Ollama + Llama 4 Scout	Free · ₹0	☐ Yes ☒ No ☐ Partial	₹0	Grounded offline engine used instead, eliminating heavy local weights.
GitHub Actions (CI/CD)	Free · ₹0	☒ Yes ☐ No ☐ Partial	₹0	Automated CI workflow executing pytest coverage and Vite build.

pytest + pytest-cov + httpx	Free · ₹0	☒ Yes ☐ No ☐ Partial	₹0	Automated QA test suite (34 tests, 94% coverage).
Faker (Python)	Free · ₹0	☒ Yes ☐ No ☐ Partial	₹0	Seeded realistic CRM synthetic dataset (660 records).
Contingency Buffer	Reserved · ₹500	☐ Yes ☒ No ☐ Partial	₹0	Reserved; zero unexpected costs incurred.

7.1 Budget Summary

Approved budget ceiling	₹2,500
Actual spend till Mid-Term (Week 9)	₹0
Actual spend, Mid-Term to Final (Weeks 10–17)	₹0
Total actual spend (full programme)	₹0 (₹1,300 estimated for production deployment)
Buffer remaining	₹2,500 (100% compliant with budget ceiling)

8. Deviations from Approved Proposal (Full Programme)

Item	Approved plan	Actual implementation	Reason for change
Model auto-selection library	PyCaret auto-select	scikit-learn compare-and-select	Carried from Mid-Term: PyCaret dependency pins conflict with Python 3.13; exact same algorithms (LogisticRegression, RandomForest, GradientBoosting, DecisionTree) evaluated via cross-validation; MLflow tracking preserved.
AI narrative generation	Gemini 2.5 Flash only	Gemini 2.5 Flash + Deterministic Grounded Fallback	Carried from Mid-Term: Protects against API rate limits and network latency during live evaluation; offline engine is factually grounded in the exact same ML brief and pipeline metrics.
E2E test framework	Playwright E2E	FastAPI TestClient + Integration Suite + Vite Build	Full automated integration suite (34 tests) and production Vite build validation cover the entire end-to-end user journey reliably across CI environments.

9. Enhancements & Additional Value-Adds

ID	Enhancement (beyond approved scope)	Why you added it / value it adds	Status	Cost impact (₹)	Evidence ID(s)
EN-01	Deterministic Grounded AI Narrative Engine	Guarantees 100% demo uptime and deterministic testability for all 3 scenarios without requiring external API keys or network access.	Done	₹0	EV-03, EV-12, EV-13
EN-02	Interactive 3-Scenario Executive Cockpit	Tab switcher on the AI Intelligence screen with color-coded operational modes and live pipeline context cards.	Done	₹0	EV-12, EV-13
EN-03	Rep Performance Quota Reference & Coaching Badges	Visual 100% quota baseline indicator and automated categorization into President's Club, On Track, and Needs Coaching.	Done	₹0	EV-10, EV-11
EN-04	Pre-Computed Score Caching on Ingestion	ML predictions (win probabilities, expected revenues, and health classifications) are pre-computed during ingestion and indexed in SQLite, reducing portal page loads to < 100ms.	Done	₹0	EV-02, EV-04

10. What Is NOT Completed + Future Scope

Pending item (be specific)	Why it wasn't completed	Flagged at Mid-Term? (Y/N)	Recommended future scope
None (100% of approved scope completed)	All approved deliverables (3 domains, 4 screens, 3 ML models, 3 AI scenarios, QA coverage ≥80%) have been fully implemented and verified.	N	Future Scope: Native bidirectional connectors for Salesforce/HubSpot via OAuth 2.0, and real-time WebSocket notifications for high-risk account alerts.

11. Risks & Blockers — Final Status

Risk / Blocker	Final Status	Mitigation taken	Final impact on delivered project
ML accuracy below target on synthetic data	Mitigated	Seeded Faker generator with realistic domain signals + cross-validation. Win Scorer AUC reached 0.9643, Rev Forecaster MAPE reached 0.0537, and Health Classifier F1 reached 0.7690.	Zero negative impact; all 3 models exceed targets.
Azure free-tier limits	Mitigated	Architected solution around free tiers (Static	Zero cost impact; 100% within

exceeded / unexpected cost		Web Apps + App Service F1); spend was ₹0 during development.	₹2,500 budget ceiling.
Gemini API rate limits during live evaluation	Mitigated	Dual-engine architecture: Gemini 2.5 Flash for live inference with seamless automatic fallback to deterministic grounded narratives.	Guarantees zero failures during live review.
FastAPI + SQLite query performance	Mitigated	Pre-computed ML predictions on ingestion and added composite indexes on hot columns.	Screen loads render in < 100ms (far beating the < 3s target).

12. Declaration & Pre-Submission Checklist

- ☒ All fields in Section 1 match my L&D Final Decision record exactly, and the Mid-Term document filename is the exact file I uploaded at Week 10.
- ☒ Section 3 lists every deliverable my approved proposal committed to across the full programme (Week 4–17), each with a D-ID and a status, reusing my Mid-Term D-IDs.
- ☒ Deliverables, evidence, and deviations carried forward from Mid-Term are clearly marked as such — nothing is silently omitted.
- ☒ Every 'Done' or 'Partial' status in Section 3 points to at least one Evidence ID in Section 4.
- ☒ Every evidence block in Section 4 has a specific caption, a Date of Development/Testing, and either a pasted full-size screenshot or a ticked carry-forward reference.
- ☒ The repository link in Section 5 is accessible to the L&D team, the stated final commit exists, and the walkthrough file paths in Section 5.1 exist in the repository.
- ☒ Section 6 coverage figures are measured (tool output attached as evidence), not estimated (94%).
- ☒ Section 7 lists every tool from my approved proposal, including ones I did not use, and reflects final total spend (₹0 actual).
- ☒ Section 8 discloses every deviation across the full programme, including ones already reported at Mid-Term.
- ☒ Section 9 (Enhancements) is filled in with cost impact per entry, and its costs reconcile with Section 7.
- ☒ Section 10 is consistent with Sections 3 and 4 and with my Mid-Term document — no contradictions.
- ☒ I have deleted all grey italic instruction text.
- ☒ I have not renamed, deleted, or reordered any section of this template.
- ☒ Document is saved as S4-I-21_Hindujashiri_FinalTermDoc.docx and uploaded to Moodle before the deadline.

Declaration: I confirm that all progress claims, evidence, costs, and coverage figures in this document are true and reflect my own individual work across the full programme.

Participant signature / name	Hindujashiri Gopu
Date of submission	03-Sep-2026 / Week 17